




# **Instruction Set Architecture ( ISA )**

# instructions

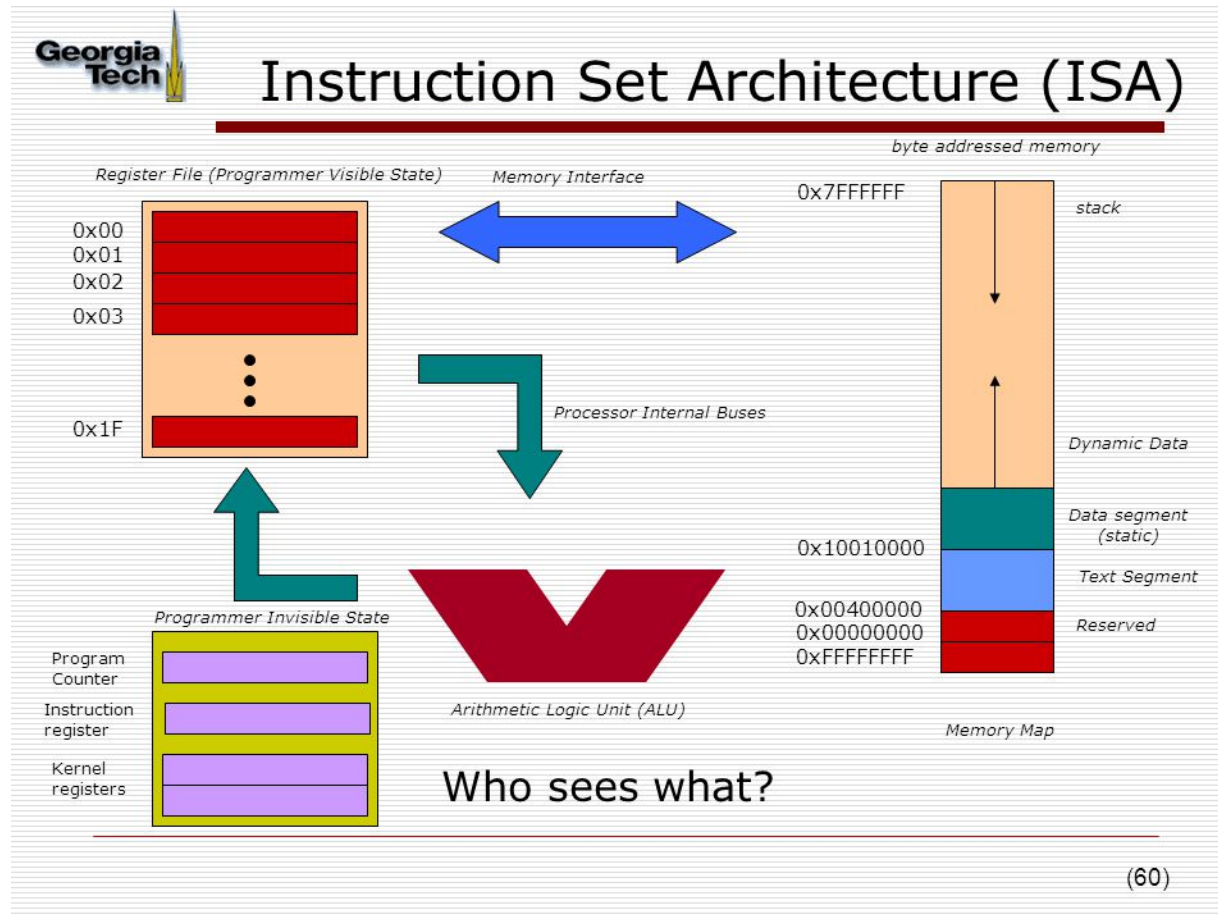




# Instruction Set Architecture

- Also called (computer) architecture
- Implementation --> actual realisation of ISA
- ISA can have multiple implementations
- ISA allows software to direct hardware
- ISA defines machine language

# ISA model





# ISA defines...

- Data types
- Memory (registers et. al.)
- Addressing modes
- Instruction set
- I/O

# What is an ISA?

- x86 --> yes!
  - Intel implements x86
  - AMD implements x86
  - Yet, both have different designs!!!
- Therefore,  
processors with different designs  
(microarchitectures) can share the same ISA

# Virtual Machines

- Java Bytecode, anyone???
  - Java programs are compiled to instruction set specific to the Java VM
  - Java VM translates bytecode to machine specific machine code
- Possible to implement one ISA on top of another one using such techniques

# How to classify ISA?

- Based on *complexity*
  - Complex Instruction Set Computer (CISC)
  - Reduced Instruction Set Computer (RISC)
- Parallelism / Word size
  - VLIW (very long instruction word)
  - LIW (long instruction word)
  - EPIC (explicitly parallel instruction computing)



# CISC ( complex )

- Single instruction can execute multiple operations (low level --> I/O, ALU, mem)
- CISC was defined *after* RISC was defined
  - But CISC came before RISC (*bizzaro!*)
  - Everything that is not RISC... is CISC



# CISC ( complex )

Designed to implement programming constructs such as:

- Procedure calls
- Loops
- Array access
- Address lookups

... into a single instruction!!!



# CISC ( complex )

CISC is awesome.....

But...

There's a catch...

Sometimes less complex instructions  
performed better...

Because programmers 'overengineered'  
(*typical!*)

# CISC

## and the return of the Superscalar

- Build superscalar implementations of CISC directly (native support)
- Advances in fast cache mediate frequent memory accesses
- Combine x86 instructions
- E.g. Pentium Pro and AMD K5
- Why didn't it catch on???

... because everyone was using x86 RISC by then

# Why did we take RISC?

- Reduced, as in, less complex than CISC
- Requires lesser cycles to execute
- Loose definition of instructions
  - simpler and smaller and general
- Has *more* (reduced set? no?) instructions

# Stanford MIPS Experiment

- Combine RISC with VLSI (very large scale integration semiconductor technology)
- 32-bit everything
  - instruction, addressing (word-addressed)
- Load/Store
  - 32 general purpose registers
- Optimising compilers!!!



# RISC

- Fixed-length instructions (mostly 32bit)
- Drawback? Code density.
- ARM, MIPS, RISC-V
  - short reduced instruction
  - instruction compression
- ARM... is the processor in your phones!  
You're all carrying a RISC in your pockets!

# What architecture is my desktop/laptop CPU?

RISC?

CISC?





# What architecture is my desktop/laptop CPU?

It implements the best of both worlds!!!





# Levels of parallelism

- Thread
- Task
- Data
- Memory
- Software



# MPP

## Massively Parallel Processing

- Large number of processors
- Simultaneously process
- Can be different computers
- e.g. Grid computing
- MPPAs – massively parallel processing arrays
- Used in supercomputers  
(this is not cluster computing)



# Grid Computing

- Distributed system
- Non-interactive workloads
- Each node (can) perform a different task
- Nodes can be heterogenous
- Nodes are not physically coupled

# GPU / GPGPU

- General Purpose computing on Graphical Processing Unit (GPGPU)
- Use GPU instead of CPU
- GPU is great at parallelisation
- Use multiple GPUs in a pipeline formation
- e.g. Nvidia CUDA
- e.g. Metal (Apple), Vulkan
- More later!!! (really later, in like, another lecture)

# Vector Processor

- Instructions operate on vectors (1D) --> SIMD !!!
- Scalar processors operate on one item
- Awesome for numerical tasks
- GPUs are kind of like vector processors
- Intel x86
  - MMX, SSE, AVX
- We'll learn about SSE soon.


# Gustafson's & Amdahl's Law

- A task executed by a system whose resources are improved compared to an initial similar system can be split into two parts:
  - a part that does not benefit from the improvement of the resources of the system;
  - a part that benefits from the improvement of the resources of the system.

# Gustafson's Law

## fixed execution time


Gustafson's Law:  $S(P) = P - a \cdot (P - 1)$





# Amdahl's Law

## fixed workload





# Parallel Programming Languages

- Actor model: Smalltalk, Erlang, Elixir
- Co-ordination: LINDA
- Dataflow: Joule
- Distributed: Julia
- Event-driven: Verilog, VHDL
- Functional: Haskell
- Logic: Prolog



# Parallel Programming Languages

- Multi-threaded: C, C++, Java
- Object-oriented: C#, Java, Smalltalk
- Message passing: Rust
- Frameworks
  - Apache Hadoop, Apache Spark, CUDA,
  - OpenCL, OpenMP



Let's discuss the assignment!!!

The assignment is *simple*.

Implement 1 **searching algorithm** and 1 **sorting algorithm** in non-parallel and parallel execution formats, and compare them. Write a report on your findings.

Submission date: FRIDAY 23:59 16 March 2018 UTC.

How to submit: blackboard, upload as many times as you want