

Vector Processors

some slides from:

Krste Asanovic

Electrical Engineering and Computer Sciences

University of California, Berkeley

Also from David Gregg

SCSS, Trinity College Dublin

The Rise of SIMD

🔑 SIMD is good for applying identical computations across many data elements

- E.g. `for(i = 0; I < 100; i++) { A[i] = B[i] + C[i];}`
- Also “data level parallelism”

🔑 SIMD is energy efficient

- Less control logic per functional unit
- Less instruction fetch and decode energy

🔑 SIMD computations tend to bandwidth-efficient and latency tolerant

- Memory systems (caches, prefetchers, etc) are good at sequential scans through arrays

🔑 Easy examples


- Dense linear algebra
- Computer graphics (which includes a lot of dense linear algebra)
- Machine vision
- Digital signal processing

🔑 Harder examples – can be made to work to gain benefits above


- Database queries
- Sorting

Vector Programming Model

Scalar Registers




Vector Registers



Vector Length Register VLR


Vector Arithmetic
Instructions

ADDV v3, v1, v2



Vector Load and
Store Instructions

LV v1, r1, r2




Vector Instruction Execution

ADDV C,A,B





*Execution using
one pipelined
functional unit*

*Execution using
four pipelined
functional units*

A[6] B[6]
A[5] B[5]
A[4] B[4]
A[3] B[3]



A[24] B[24] A[25] B[25] A[26] B[26] A[27] B[27]
A[20] B[20] A[21] B[21] A[22] B[22] A[23] B[23]
A[16] B[16] A[17] B[17] A[18] B[18] A[19] B[19]
A[12] B[12] A[13] B[13] A[14] B[14] A[15] B[15]



Vector Memory-Memory versus Vector Register Machines

- Vector memory-memory instructions hold all vector operands in main memory (not vector registers)
- The first vector machines, CDC Star-100 ('73) and TI ASC ('71), were memory-memory machines
- Cray-1 ('76) was first vector register machine

Example Source Code
for (i=0; i<N; i++)
{
 C[i] = A[i] + B[i];
 D[i] = A[i] - B[i];
}


Vector Memory-Memory Code

```
ADDV C, A, B  
SUBV D, A, B
```


Vector Register Code

```
LV V1, A  
LV V2, B  
ADDV V3, V1, V2  
SV V3, C  
SUBV V4, V1, V2  
SV V4, D
```

Vector Unit Structure



Multimedia Extensions (aka SIMD extensions)



☞ Very short vectors added to existing ISAs for microprocessors

☞ Use existing 64-bit registers split into 2x32b or 4x16b or 8x8b


- Lincoln Labs TX-2 from 1957 had 36b datapath split into 2x18b or 4x9b

- Newer designs have wider registers

- 128b for PowerPC AltiVec, Intel SSE2/3/4

- 256b for Intel AVX

☞ Single instruction operates on all elements within register



Multimedia Extensions versus Vectors

Limited instruction set:

- no vector length control
- no strided load/store or scatter/gather
- unit-stride loads must be aligned to 64/128-bit boundary

Limited vector register length:

- requires superscalar dispatch to keep multiply/add/load units busy
- loop unrolling to hide latencies increases register pressure

Trend towards fuller vector support in microprocessors

- Better support for misaligned memory accesses
- Support of double-precision (64-bit floating-point)
- Intel AVX spec, 256b vector registers (expandable up to 1024b)

Modern Vector Computers

- Modern vector machines have relatively short vectors
 - ARM NEON: 16 bytes, Intel SSE: 16 bytes, Intel AVX:32 bytes
 - But the trend is for them to grow longer
 - AVX-512 has 64 byte vectors, i.e. 16 floats
- Older vector machines used much longer vectors
 - Cray 1 vector supercomputer used vectors of 64 floats, i.e. 256 bytes
- Modern vector machines always use vector registers for everything
 - Modern short vector registers of 16 or 32 bytes don't require much space on chip
- Some (but not all) older vector machines took data directly from memory

Modern Vector Computers

- Modern vector machines use a separate arithmetic/floating point unit per lane
 - Four parallel floating point adders/multipliers in SSE implementations
- Older vector machines used as little as one arithmetic/floating point unit to implement vector instruction
 - But very deeply pipelined
 - Goal was to push as much work through the pipelined FP unit as possible
- Vector architectures are increasingly important
 - Especially for low-energy computation
 - It's worthwhile looking back to the time when vector computers were last really popular and successful

Supercomputers

Definition of a supercomputer:

- Fastest machine in world at given task
- A device to turn a compute-bound problem into an I/O bound problem
- Any machine costing \$30M+
- Any machine designed by Seymour Cray

CDC6600 (Cray, 1964) regarded as first supercomputer

Supercomputer Applications

Typical application areas

- Military research (nuclear weapons, cryptography)
- Scientific research
- Weather forecasting
- Oil exploration
- Industrial design (car crash simulation)

All involve huge computations on large data sets

In 70s-80s, Supercomputer \equiv Vector Machine

Vector Supercomputers

Epitomized by Cray-1, 1976:

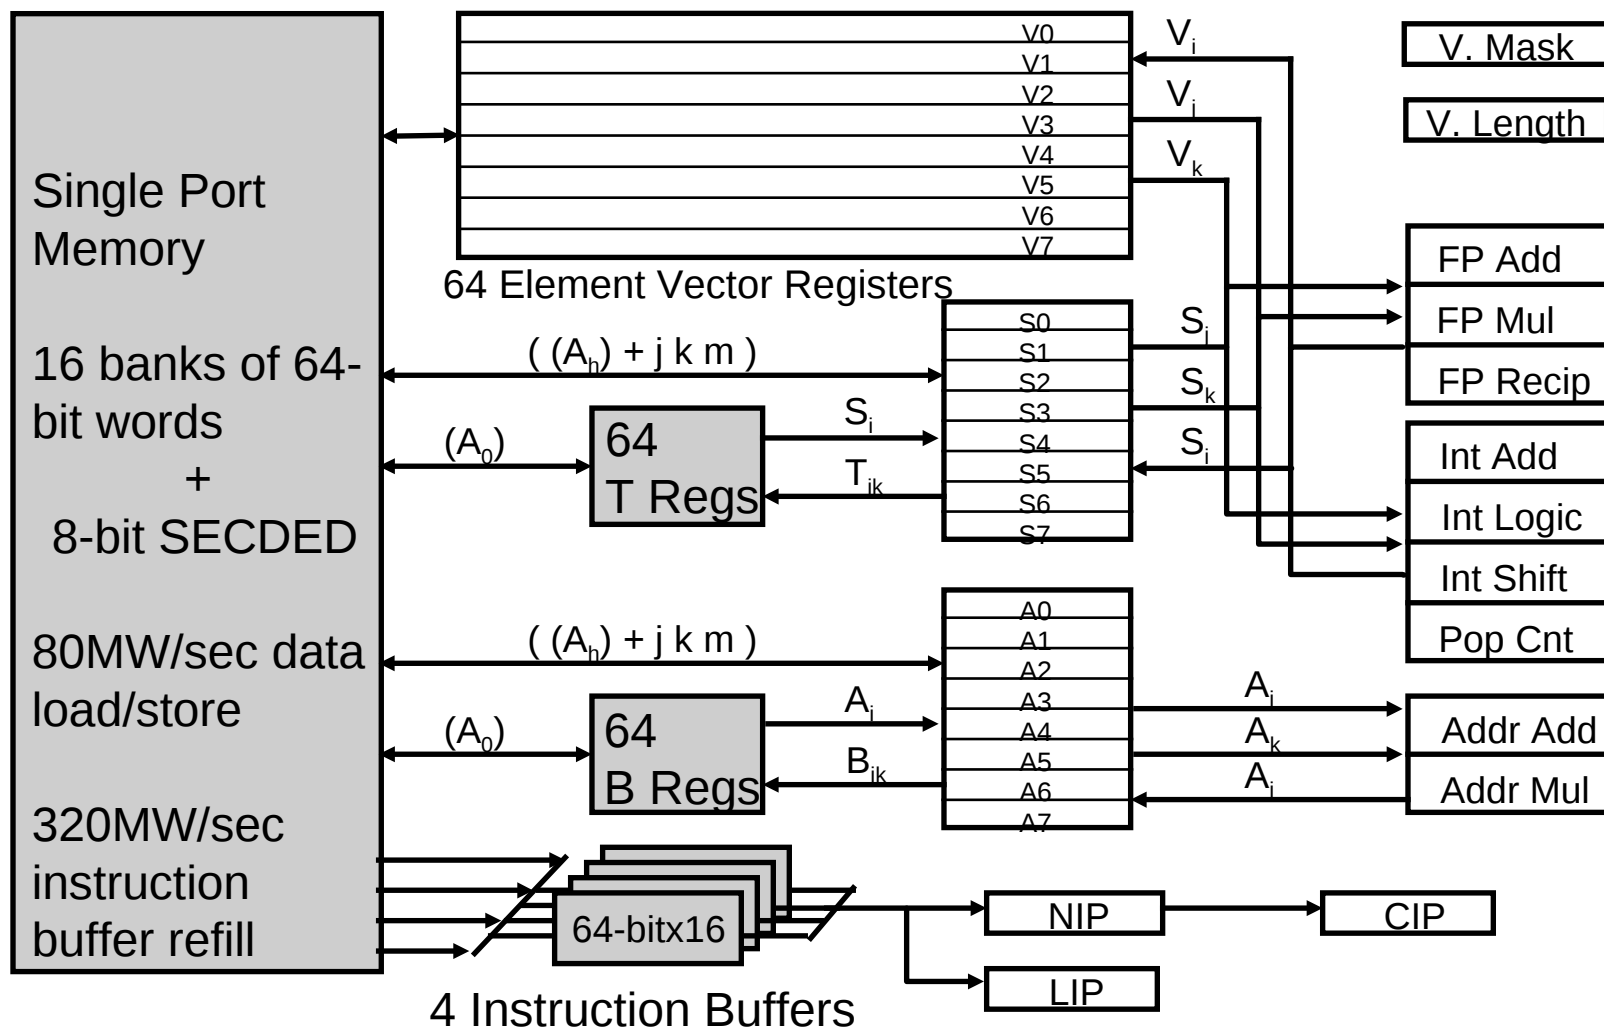
Scalar Unit + Vector Extensions

- Load/Store Architecture
- Vector Registers
- Vector Instructions
- Hardwired Control
- Highly Pipelined Functional Units
- Interleaved Memory System
- No Data Caches
- No Virtual Memory

Cray-1 (1976)




Cray-1 (1976)



memory bank cycle 50 ns processor cycle 12.5 ns (80MHz)

Vector Chaining

- Vector version of register bypassing
 - introduced with Cray-1




Vector Chaining Advantage

- Without chaining, must wait for last element of result to be written before starting dependent instruction




- With chaining, can start dependent instruction as soon as first result appears



Vector Instruction Parallelism

Can overlap execution of multiple vector instructions

- example machine has 32 elements per vector register and 8 lanes




Complete 24 operations/cycle while issuing 1 short instruction/cycle



Vector Startup

Two components of vector startup penalty

- functional unit latency (time through pipeline)
- dead time or recovery time (time before another vector instruction can start down pipeline)



Dead Time and Short Vectors



T0, Eight lanes

No dead time

100% efficiency with 8 element vectors

Vector Scatter/Gather

Want to vectorize loops with indirect accesses:

```
for (i=0; i<N; i++)  
    A[i] = B[i] + C[D[i]]
```

Indexed load instruction (*Gather*)

```
LV vD, rD          # Load indices in D vector  
LVI vC, rC, vD      # Load indirect from rC base  
LV vB, rB           # Load B vector  
ADDV.D vA, vB, vC   # Do add  
SV vA, rA           # Store result
```

Vector Scatter/Gather

Scatter example:

```
for (i=0; i<N; i++)  
    A[B[i]]++;
```

Vector Conditional Execution

Problem: Want to vectorize loops with conditional code:

```
for (i=0; i<N; i++)  
    if (A[i]>0) then  
        A[i] = B[i];
```

Solution: Add vector *mask* (or *flag*) registers

- vector version of predicate registers, 1 bit per element

...and *maskable* vector instructions

- vector operation becomes NOP at elements where mask bit is clear

Code example:


```
CVM                # Turn on all elements  
LV vA, rA           # Load entire A vector  
SGTVS.D vA, F0      # Set bits in mask register where A>0  
LV vA, rB           # Load B vector into A under mask  
SV vA, rA           # Store A back to memory under mask
```

Masked Vector Instructions

Simple Implementation

**execute all N operations, turn off result
writeback according to mask**

M[7]=1	A[7]	B[7]
M[6]=0	A[6]	B[6]
M[5]=1	A[5]	B[5]
M[4]=1	A[4]	B[4]
M[3]=0	A[3]	B[3]




M[2]=0
M[1]=1
M[0]=0

Write Enable *Write data port*

Density-Time Implementation

**scan mask vector and only execute
elements with non-zero masks**

M[7]=1	A[7]	B[7]
M[6]=0		
M[5]=1		
M[4]=1		
M[3]=0		
M[2]=0		
M[1]=1		
M[0]=0		



Write data port

Vector Reductions

Problem: Loop-carried dependence on reduction variables

```
sum = 0;
for (i=0; i<N; i++)
    sum += A[i]; # Loop-carried dependence on sum
```

Solution: Re-associate operations if possible, use binary tree to perform reduction

Rearrange as:

```
sum[0:VL-1] = 0 # Vector of VL partial sums
for(i=0; i<N; i+=VL) # Stripmine VL-sized chunks
    sum[0:VL-1] += A[i:i+VL-1]; # Vector sum
# Now have VL partial sums in one vector register
do {
    VL = VL/2; # Halve vector length
    sum[0:VL-1] += sum[VL:2*VL-1] # Halve no. of partials
} while (VL>1)
```