

---

# CS3014: Concurrent Systems

## Static & Dynamic Instruction Scheduling

Slides originally developed by  
Drew Hilton, Amir Roth, Milo Martin and Joe Devietti  
at University of Pennsylvania

---

# **Instruction Scheduling & Limitations**

# Instruction Scheduling

---

- Scheduling: act of finding independent instructions
  - “Static” done at compile time by the compiler (software)
  - “Dynamic” done at runtime by the processor (hardware)
- Why schedule code?
  - Scalar pipelines: fill in load-to-use delay slots to improve CPI
  - Superscalar: place independent instructions together
    - As above, load-to-use delay slots
    - Allow multiple-issue decode logic to let them execute at the same time

---

# **Dynamic (Execution-time) Instruction Scheduling**

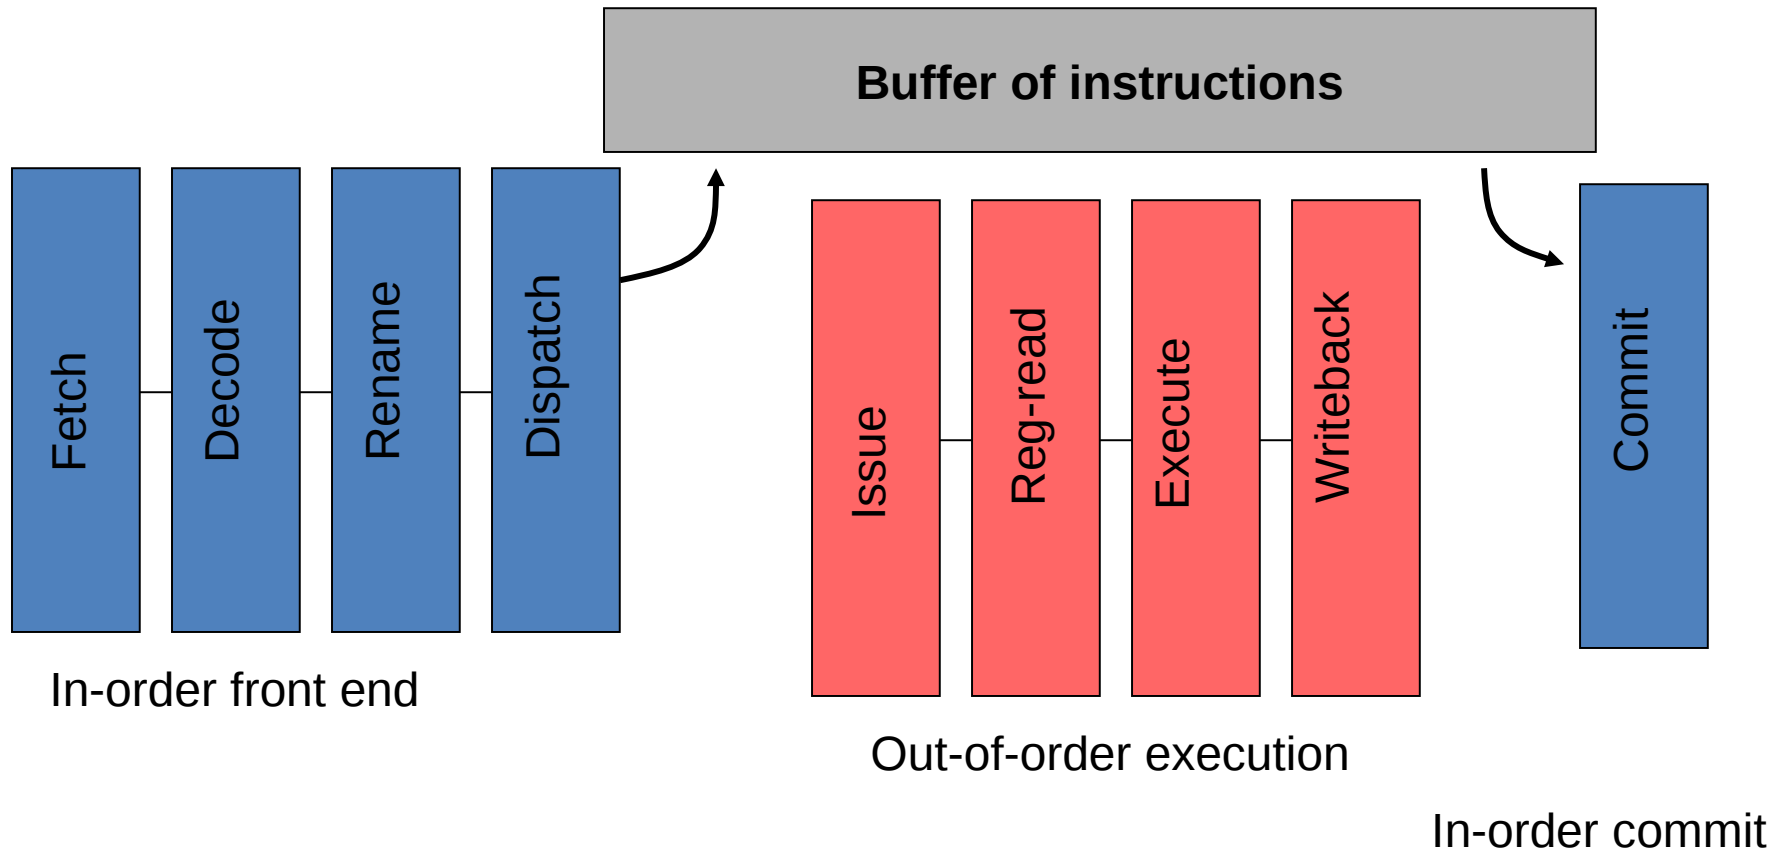
# Can Hardware Overcome These Limits?

---

- **Dynamically-scheduled processors**
  - Also called “out-of-order” processors
  - Hardware re-schedules instructions...
  - ...within a sliding window of instructions
  - As with pipelining and superscalar, ISA unchanged
    - Same hardware/software interface, appearance of in-order
- Increases scheduling scope
  - Does loop unrolling transparently!
  - Uses branch prediction to “unroll” branches
- Examples:
  - Pentium Pro/II/III (3-wide), Core 2 (4-wide), Alpha 21264 (4-wide), MIPS R10000 (4-wide), Power5 (5-wide)

# Out-of-Order Pipeline

---



# Out-of-Order Execution

---

- Also called “Dynamic scheduling”
  - Done by the hardware on-the-fly during execution
- Looks at a “window” of instructions waiting to execute
  - Each cycle, picks the next ready instruction(s)
- Two steps to enable out-of-order execution:  
Step #1: Register renaming – to avoid “false” dependencies  
Step #2: Dynamically schedule – to enforce “true” dependencies
- Key to understanding out-of-order execution:
  - **Data dependencies**

# Dependence types

---

- **RAW** (Read After Write) = “true dependence” (true)

mul r0 \* r1 → **r2**

...

add **r2** + r3 → r4

- **WAW** (Write After Write) = “output dependence” (false)

mul r0 \* r1 → **r2**

...

add r1 + r3 → **r2**

- **WAR** (Write After Read) = “anti-dependence” (false)

mul r0 \* **r1** → r2

...

add r3 + r4 → **r1**

- WAW & WAR are “false”, Can be **totally eliminated** by “renaming”



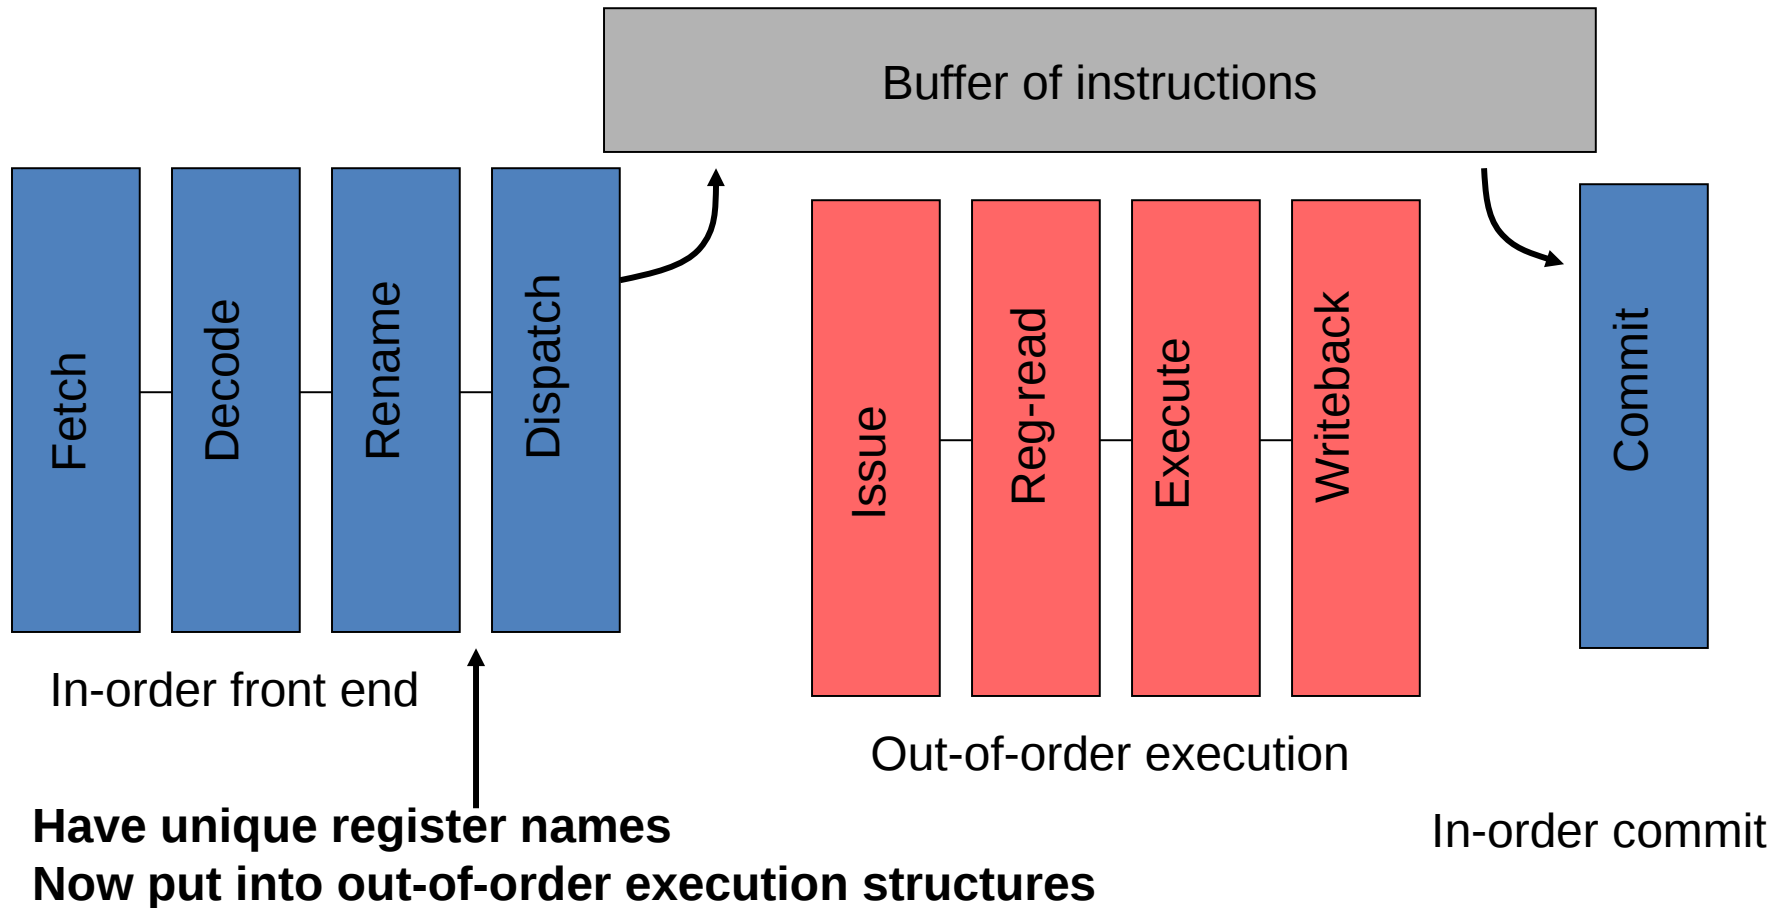
# Step #1: Register Renaming

- To eliminate register conflicts/hazards
- “Architected” vs “Physical” registers – level of indirection
  - Names: **r1, r2, r3**
  - Locations: **p1, p2, p3, p4, p5, p6, p7**
  - Original mapping: **r1→p1, r2→p2, r3→p3**, p4–p7 are “available”


Time	MapTable			FreeList	Original insns	Renamed insns
	r1	r2	r3	p4, p5, p6, p7	add r2, r3→r1	add p2, p3→ <b>p4</b>
	p1	p2	p3	p5, p6, p7	sub r2, r1→r3	sub p2, <b>p4</b> → <b>p5</b>
	p4	p2	p3	p6, p7	mul r2, r3→r3	mul p2, <b>p5</b> → <b>p6</b>
	p4	p2	p5	p7	div r1, #4→r1	div <b>p4</b> , #4→ <b>p7</b>
	p4	p2	p6			

- Renaming – conceptually write each register once
  - + Removes **false** dependences
  - + Leaves **true** dependences intact!
- When to reuse a physical register? After overwriting instruction is complete

# Out-of-order Pipeline



## Step #2: Dynamic Scheduling



Ready Table

	P2	P3	P4	P5	P6	P7
Time	Yes	Yes				
	Yes	Yes	Yes			
	Yes	Yes	Yes	Yes		Yes
	Yes	Yes	Yes	Yes	Yes	Yes

add p2, p3→p4

sub p2, p4→p5


mul p2, p5→p6

and div p4, 4→p7

- Instructions fetch/decoded/renamed into *Instruction Buffer*
  - Also called “instruction window” or “instruction scheduler”
- Instructions (conceptually) check ready bits every cycle
  - Execute oldest “ready” instruction, set output as “ready”

# Dynamic Scheduling/Issue Algorithm

---

- Data structures:
  - Ready table[phys\_reg]  yes/no (part of “issue queue”)
- Algorithm at “issue” stage (prior to read registers):  
foreach instruction:  
if table[insn.phys\_input1] == ready &&  
table[insn.phys\_input2] == ready then  
    insn is “ready”  
select the oldest “ready” instruction  
table[insn.phys\_output] = ready
- Multiple-cycle instructions? (such as loads)
  - For an instruction with latency of N, set “ready” bit N-1 cycles in future

---

# Register Renaming

# Register Renaming Algorithm (Simplified)

---

- Two key data structures:
  - `maptable[architectural_reg] → physical_reg`
  - Free list: allocate (new) & free registers (implemented as a queue)
    - ignore freeing of registers for now
- Algorithm: at “decode” stage for each instruction:
  - Rewrites instruction with “physical” registers (rather than “architectural” registers)

```
insn.phys_input1 = maptable[insn.arch_input1]  
insn.phys_input2 = maptable[insn.arch_input2]
```

```
new_reg = new_phys_reg()  
maptable[insn.arch_output] = new_reg  
insn.phys_output = new_reg
```

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Renaming example

---

xor **r1** ^ **r2** → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

————→ xor **p1** ^ **p2** →

<b>r1</b>	<b>p1</b>
<b>r2</b>	<b>p2</b>
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list



# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

————→ xor p1 ^ p2 → **p6**

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

<b>p6</b>
p7
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → **r3**  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

————→ xor p1 ^ p2 → p6

r1	p1
r2	p2
<b>r3</b>	<b>p6</b>
r4	p4
r5	p5

Map table


p7
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add **r3** + **r4** → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add **p6** + **p4** →

r1	p1
r2	p2
<b>r3</b>	<b>p6</b>
<b>r4</b>	<b>p4</b>
r5	p5

Map table


p7
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → **p7**

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table

<b>p7</b>
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → **r4**  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7

r1	p1
r2	p2
r3	p6
<b>r4</b>	<b>p7</b>
r5	p5

Map table


p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub **r5** - **r2** → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub **p5** - **p2** →

r1	p1
<b>r2</b>	<b>p2</b>
r3	p6
r4	p7
<b>r5</b>	<b>p5</b>

Map table


p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → **p8**

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table


<b>p8</b>
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → **r3**  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8

r1	p1
r2	p2
<b>r3</b>	<b>p8</b>
r4	p7
r5	p5

Map table

p9
p10

Free-list




# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi **r3** + 1 → r1

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi **p8** + 1 →



r1	p1
r2	p2
<b>r3</b>	<b>p8</b>
r4	p7
r5	p5

Map table

p9
p10


Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → **p9**



r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

<b>p9</b>
p10

Free-list

# Renaming example


---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → **r1**

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

<b>r1</b>	<b>p9</b>
r2	p2
r3	p8
r4	p7
r5	p5

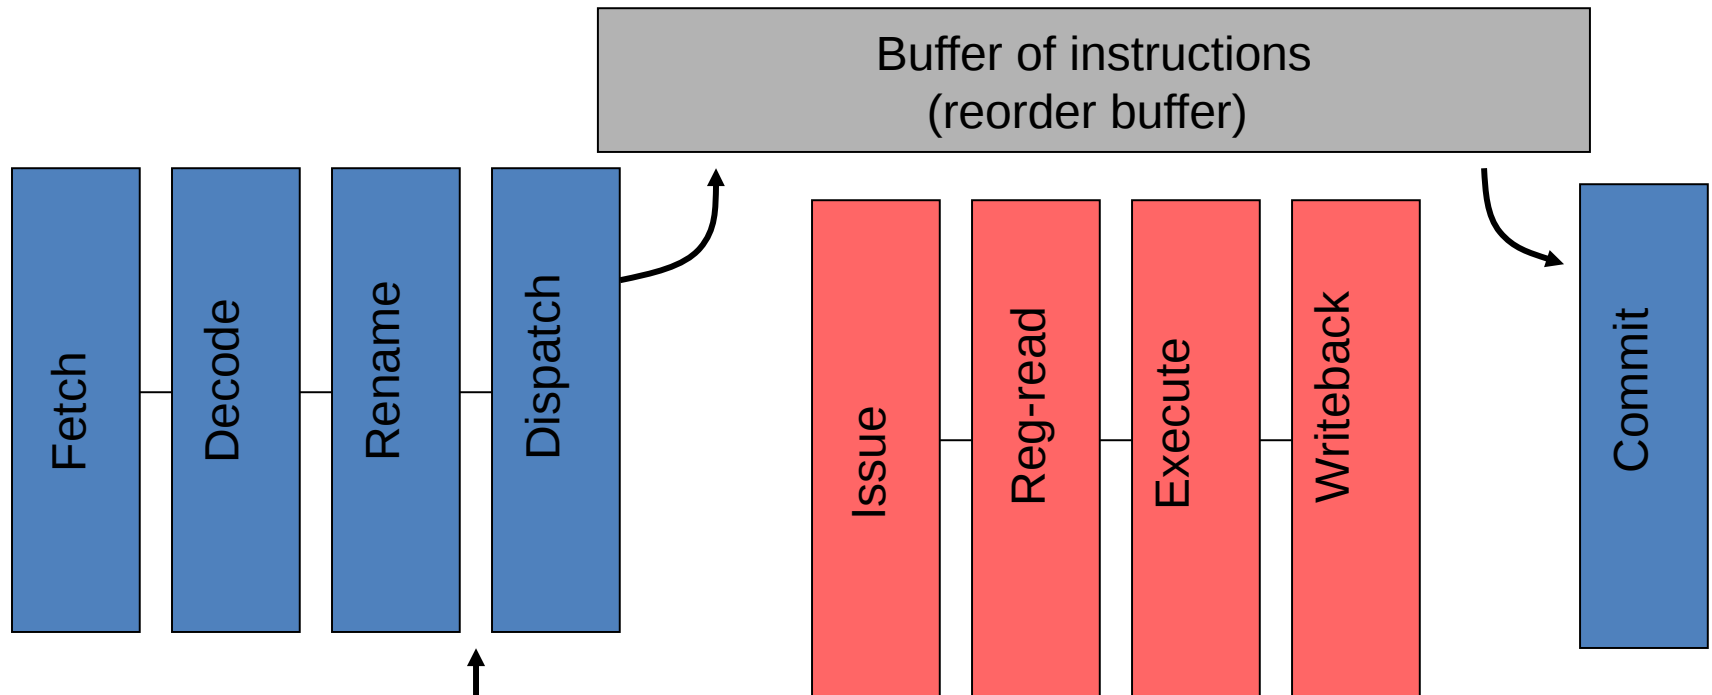
Map table



Free-list

# Out-of-order Pipeline

---



Have unique register names  
Now put into out-of-order execution structures

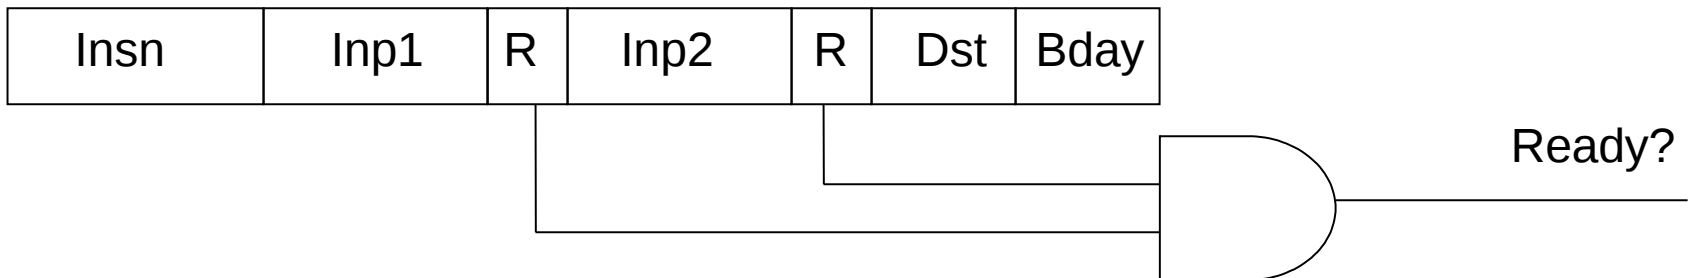
---

# **Dynamic Instruction Scheduling Mechanisms**

# Dispatch

---

- Put renamed instructions into out-of-order structures
- Re-order buffer (ROB)
  - Holds instructions from Fetch through Commit
- Issue Queue
  - Central piece of scheduling logic
  - Holds instructions from Dispatch through Issue
  - Tracks ready inputs
    - Physical register names + ready bit
    - “AND” the bits to tell if ready



# Dispatch Steps

---

- Allocate Issue Queue (IQ) slot
  - Full? Stall
- Read **ready bits** of inputs
  - 1-bit per physical reg
- Clear **ready bit** of output in table
  - Instruction has not produced value yet
- Write data into Issue Queue (IQ) slot

# Dispatch Example

---

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

**Issue Queue**

Insn	Inp1	R	Inp2	R	Dst	Bday

**Ready bits**

p1	y
p2	y
p3	y
p4	y
p5	y
p6	y
p7	y
p8	y
p9	y



# Dispatch Example

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	y
p8	y
p9	y

# Dispatch Example

---

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
<b>p7</b>	<b>n</b>
p8	y
p9	y

# Dispatch Example

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1
sub	p5	y	p2	y	p8	2

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
p8	n
p9	y

# Dispatch Example

---

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

**Issue Queue**

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1
sub	p5	y	p2	y	p8	2
addi	p8	n	---	y	p9	3


**Ready bits**

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
p8	n
p9	n

# Out-of-order pipeline


---

- Execution (out-of-order) stages
- **Select** ready instructions
  - Send for execution
- **Wakeup** dependents



# Dynamic Scheduling/Issue Algorithm

---

- Data structures:
  - Ready table[phys\_reg]  yes/no (part of issue queue)
- Algorithm at “schedule” stage (prior to read registers):  
foreach instruction:  
if table[insn.phys\_input1] == ready &&  
table[insn.phys\_input2] == ready then  
    insn is “ready”  
select the oldest “ready” instruction  
table[insn.phys\_output] = ready

# Issue = Select + Wakeup

---

- **Select** oldest of “ready” instructions
  - “xor” is the oldest ready instruction below
  - “xor” and “sub” are the two oldest ready instructions below
  - ✂ Note: may have resource constraints: i.e. load/store/floating point

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1
sub	p5	y	p2	y	p8	2
addi	p8	n	---	y	p9	3

**Ready!**

**Ready!**

# Issue = Select + Wakeup

- Wakeup dependent instructions
  - Search for destination (Dst) in inputs & set “ready” bit
    - Implemented with a special memory array circuit called a Content Addressable Memory (CAM)
  - Also update ready-bit table for future instructions

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0
add	p6	y	p4	y	p7	1
sub	p5	y	p2	y	p8	2
addi	p8	y	---	y	p9	3

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	y
p7	n
p8	y
p9	n

- For multi-cycle operations (loads, floating point)
  - Wakeup deferred a few cycles
  - Include checks to avoid structural hazards



# Note: Content Addressable Memory

---

- A content addressable memory (CAM) is indexed by the **content** of each location, not by the address
  - Sometimes known as associative memory
- It compares an input “key” against a table of keys, and returns the location of the key in the table
  - In software this might be implemented with a hash table
  - Hardware hash table is also possible, but potentially slow
- To search all locations in a single cycle
  - You need to be able to compare the search key to all keys in the table simultaneously
    - This requires a *\*lot\** of hardware
    - Fast CAMs are very hardware expensive
  - If you need to be able to do multiple searches in the same cycle, the hardware requirements are even greater

# Issue


---

- **Select/Wakeup** one cycle
- Dependent instructions execute on back-to-back cycles
  - Next cycle: add/addi are ready:

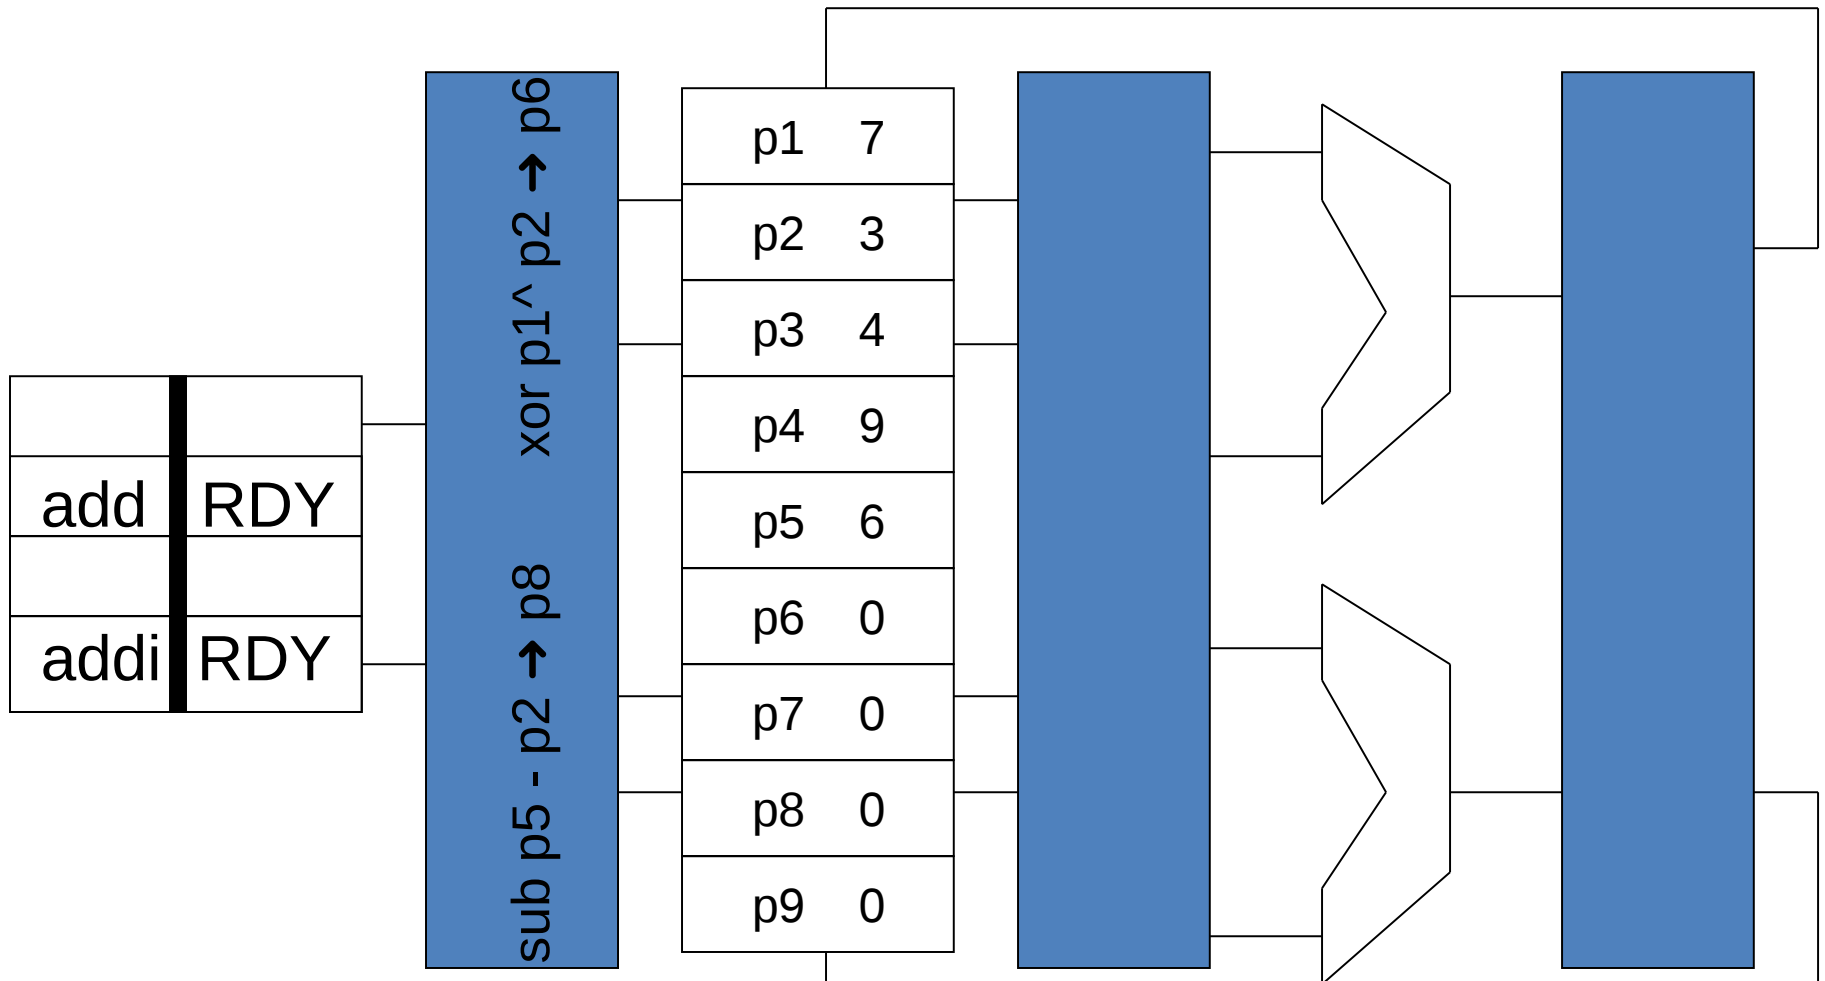
Insn	Inp1	R	Inp2	R	Dst	Bday
add	p6	y	p4	y	p7	1
addi	p8	y	---	y	p9	3

- Issued instructions are removed from issue queue
  - Free up space for subsequent instructions

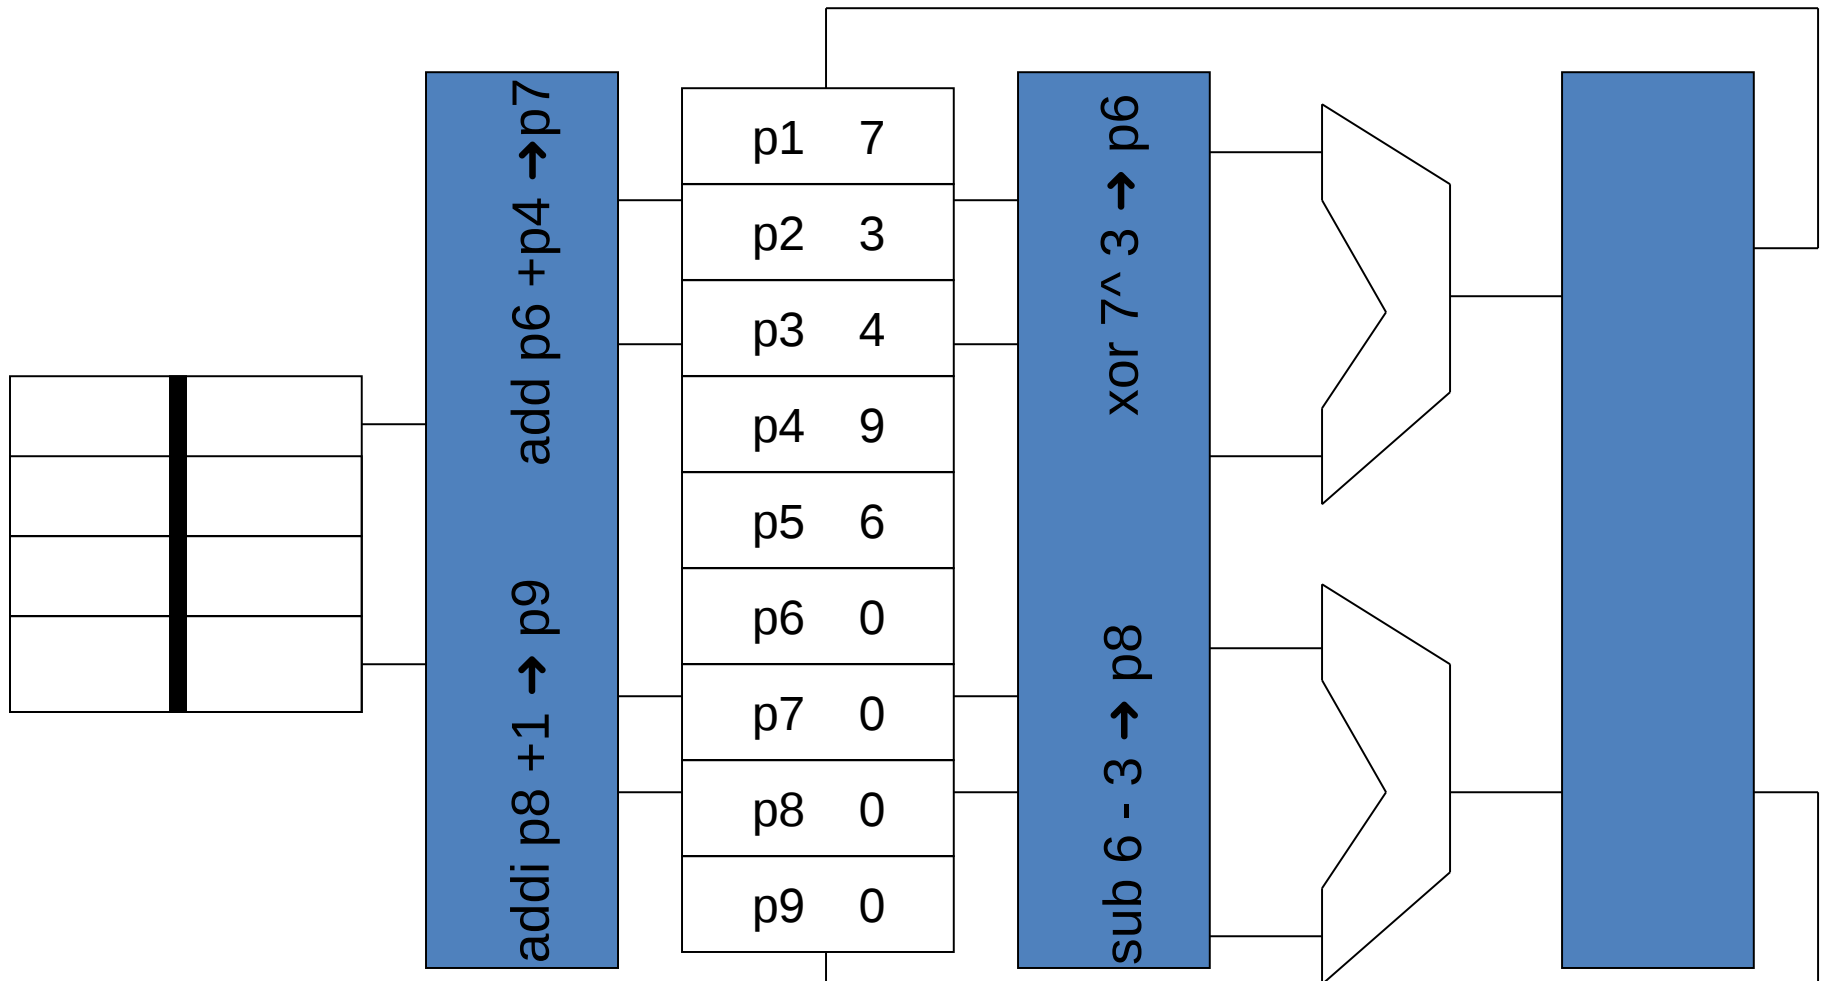
# OOO execution (2-wide)



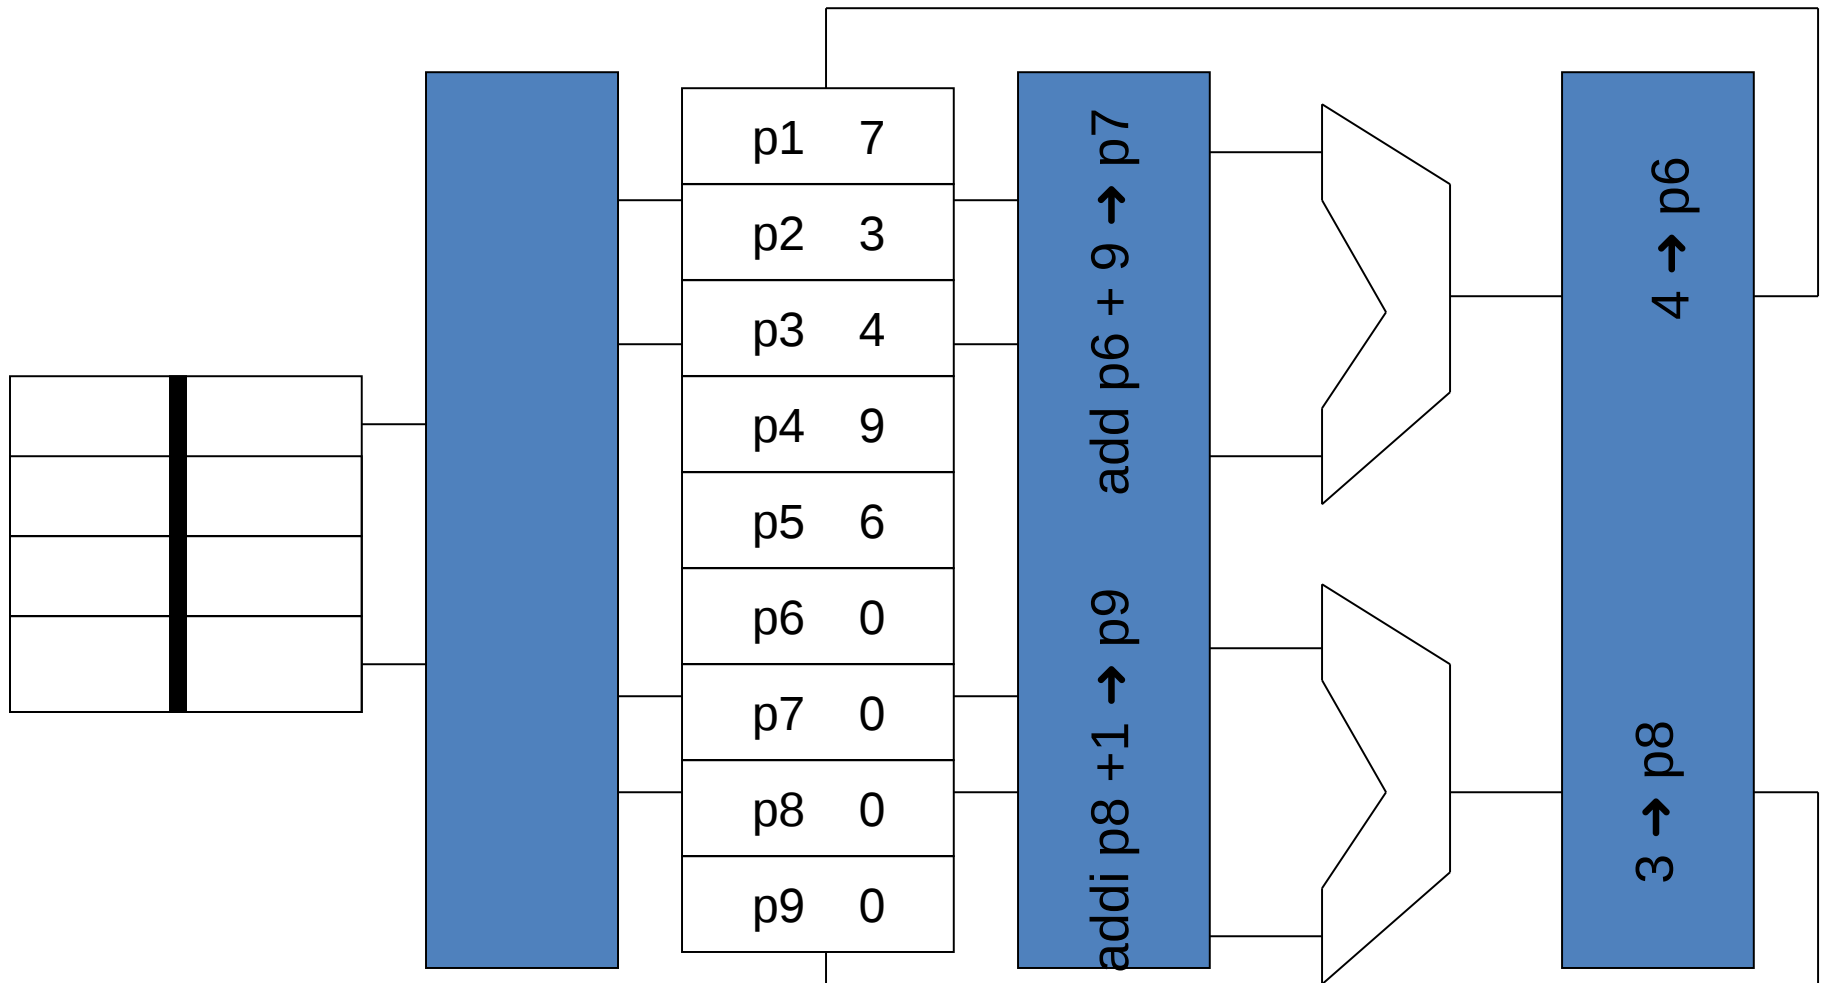
# OOO execution (2-wide)




# OOO execution (2-wide)




# OOO execution (2-wide)



# OOO execution (2-wide)




# OOO execution (2-wide)





# OOO execution (2-wide)

**Note similarity  
to in-order**



---

# Out-of-Order: Benefits & Challenges








# Dynamic Scheduling Operation (Recap)

---

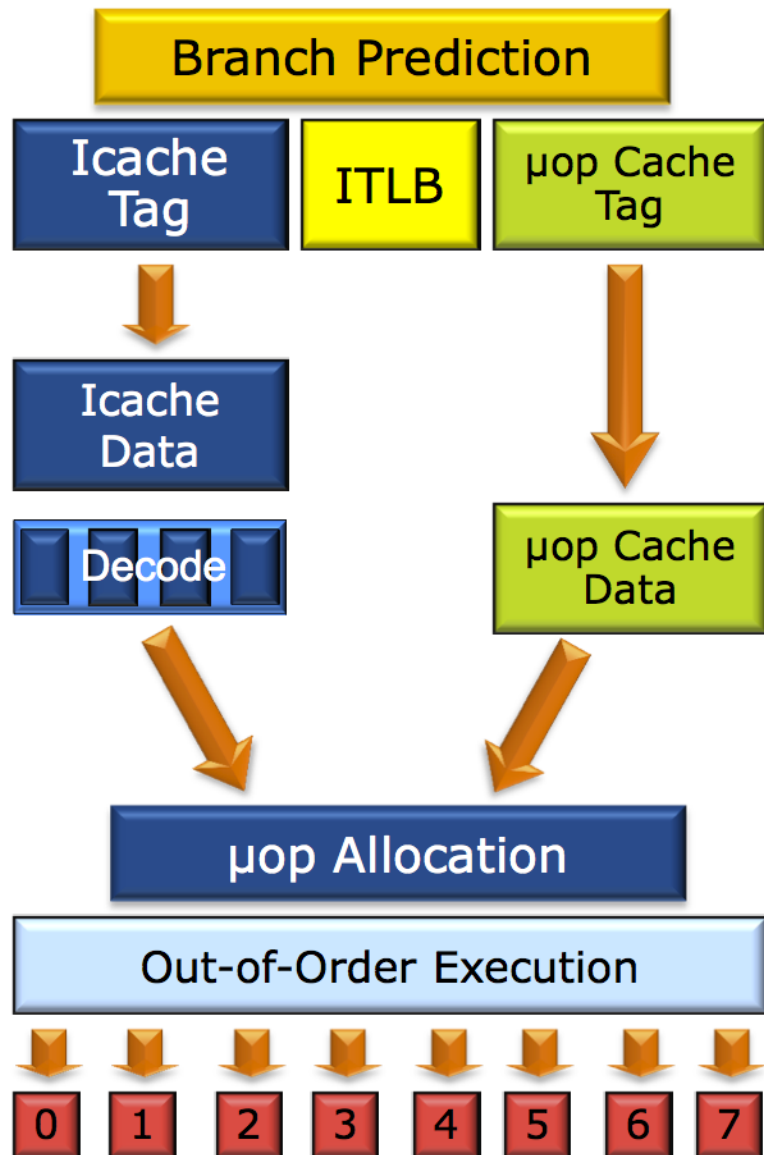
- Dynamic scheduling
  - Totally in the hardware (not visible to software)
  - Also called “out-of-order execution” (OoO)
- Fetch many instructions into instruction window
  - Use branch prediction to speculate past (multiple) branches
  - Flush pipeline on branch misprediction
- Rename registers to avoid false dependencies
- Execute instructions as soon as possible
  - Register dependencies are known
  - Handling memory dependencies is harder
- “Commit” instructions in order
  - Anything strange happens before commit, just flush the pipeline

# Haswell Buffer Sizes

Extract more parallelism in every generation

	Nehalem	Sandy Bridge	Haswell	
Out-of-order Window	128	168	192	
In-flight Loads	48	64	72	
In-flight Stores	32	36	42	
Scheduler Entries	36	54	60	
Integer Register File	N/A	160	168	
FP Register File	N/A	144	168	
Allocation Queue	28/thread	28/thread	56	

# Haswell Core at a Glance



## Next generation branch prediction

- Improves performance *and* saves wasted work

## Improved front-end

- Initiate TLB and cache misses speculatively
- Handle cache misses in parallel to hide latency
- Leverages improved branch prediction

## Deeper buffers

- Extract more instruction parallelism
- More resources when running a single thread

## More execution units, shorter latencies

- Power down when not in use

## More load/store bandwidth

- Better prefetching, better cache line split latency & throughput, double L2 bandwidth
- New modes save power without losing performance

## No pipeline growth

- Same branch misprediction latency
- Same L1/L2 cache latency

# Core Cache Size/Latency/Bandwidth

Metric	Nehalem	Sandy Bridge	Haswell
L1 Instruction Cache	32K, 4-way	32K, 8-way	32K, 8-way
L1 Data Cache	32K, 8-way	32K, 8-way	32K, 8-way
Fastest Load-to-use	4 cycles	4 cycles	4 cycles
Load bandwidth	16 Bytes/cycle	32 Bytes/cycle (banked)	64 Bytes/cycle
Store bandwidth	16 Bytes/cycle	16 Bytes/cycle	32 Bytes/cycle
L2 Unified Cache	256K, 8-way	256K, 8-way	256K, 8-way
Fastest load-to-use	10 cycles	11 cycles	11 cycles
Bandwidth to L1	32 Bytes/cycle	32 Bytes/cycle	64 Bytes/cycle
L1 Instruction TLB	4K: 128, 4-way 2M/4M: 7/thread	4K: 128, 4-way 2M/4M: 8/thread	4K: 128, 4-way 2M/4M: 8/thread
L1 Data TLB	4K: 64, 4-way 2M/4M: 32, 4-way 1G: fractured	4K: 64, 4-way 2M/4M: 32, 4-way 1G: 4, 4-way	4K: 64, 4-way 2M/4M: 32, 4-way 1G: 4, 4-way
L2 Unified TLB	4K: 512, 4-way	4K: 512, 4-way	4K+2M shared: 1024, 8-way

All caches use 64-byte lines

# OoO Execution is all around us

---

- Example phone CPU: Qualcomm Krait 400 processor
  - based on ARM Cortex A15 processor
  - **out-of-order** 2.5GHz quad-core
  - 3-wide fetch/decode
  - 4-wide issue
  - 11-stage integer pipeline
  - 28nm process technology
  - 4/4KB DM L1\$, 16/16KB 4-way SA L2\$, 2MB 8-way SA L3\$

# Out of Order: Benefits

---

- Allows speculative re-ordering
  - Loads / stores
  - Branch prediction to look past branches
- Done by hardware
  - Compiler may want different schedule for different hw configs
  - Hardware has only its own configuration to deal with
- Schedule can change due to cache misses
- **Memory-level parallelism**
  - Executes “around” cache misses to find independent instructions
  - Finds and initiates independent misses, reducing memory latency
    - Especially good at hiding L2 hits (~12 cycles in Core i7)



# Challenges for Out-of-Order Cores

---

- Design complexity
  - More complicated than in-order? Certainly!
  - But, we have managed to overcome the design complexity
- Clock frequency
  - Can we build a “high ILP” machine at high clock frequency?
  - Yep, with some additional pipe stages, clever design
- Limits to (efficiently) scaling the window and ILP
  - Large physical register file
  - Fast register renaming/wakeup/select/load queue/store queue
    - Active areas of micro-architectural research
  - Branch & memory depend. prediction (limits effective window size)
    - 95% branch mis-prediction: 1 in 20 branches, or 1 in 100 insn.
  - Plus all the issues of building “wide” in-order superscalar
- Power efficiency
  - Today, even mobile phone chips are out-of-order cores

# Redux: HW vs. SW Scheduling

---

- Static scheduling
  - Performed by compiler, limited in several ways
- Dynamic scheduling
  - Performed by the hardware, overcomes limitations
- Static limitation → dynamic mitigation
  - Number of registers in the ISA → register renaming
  - Scheduling scope → branch prediction & speculation
  - Inexact memory aliasing information → speculative memory ops
  - Unknown latencies of cache misses → execute when ready
- Which to do? **Compiler does what it can, hardware the rest**
  - Why? dynamic scheduling needed to sustain more than 2-way issue
  - **Helps with hiding memory latency** (execute around misses)
  - Intel Core i7 is four-wide execute w/ scheduling window of 100+
  - Even mobile phones have dynamically scheduled cores (ARM A9, A15)